

# Job Output

```
rs6785@login-2-1:~  
[rs6785@login-2-1 ~]$ hadoop jar /opt/cloudera/parcels/CDH/lib/hadoop-mapreduce/  
hadoop-streaming.jar -D mapreduce.job.reduces=1 -files hdfs://dumbo/user/rs6785/  
class1/python_code/MaxTemperatureMapper.py,hdfs://dumbo/user/rs6785/class1/pytho  
n_code/MaxTemperatureReducer.py -mapper "python MaxTemperatureMapper.py" -reduce  
r "python MaxTemperatureReducer.py" -input /user/rs6785/class1/temperatureInputs  
.txt -output /user/rs6785/class1/output  
packageJobJar: [ ] [/opt/cloudera/parcels/CDH-5.15.2-1.cdhs.15.2.p0.3/jars/hadoop  
-streaming-2.6.0-cdh5.15.2.jar] /tmp/streamjob807489575681978287.jar tmpDir=null  
19/09/09 12:11:07 INFO mapreduce.JobSubmissionFiles: Permissions on staging dire  
ctory /user/rs6785/.staging are incorrect: rwxrwx---. Fixing permissions to corr  
ect value rwx-----  
19/09/09 12:11:07 INFO mapred.FileInputFormat: Total input paths to process : 1  
19/09/09 12:11:09 INFO mapreduce.JobSubmitter: number of splits:2  
19/09/09 12:11:08 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_15  
64775008355_5017  
19/09/09 12:11:08 INFO impl.YarnClientImpl: Submitted application application_15  
64775008355_5017  
19/09/09 12:11:08 INFO mapreduce.Job: The url to track the job: http://babar.es.  
its.nyu.edu:8088/proxy/application_1564775008355_5017/  
19/09/09 12:11:08 INFO mapreduce.Job: Running job: job_1564775008355_5017  
19/09/09 12:11:13 INFO mapreduce.Job: Job job_1564775008355_5017 running in uber  
mode : false  
19/09/09 12:11:13 INFO mapreduce.Job: map 0% reduce 0%  
19/09/09 12:11:17 INFO mapreduce.Job: map 50% reduce 0%  
19/09/09 12:11:18 INFO mapreduce.Job: map 100% reduce 0%  
19/09/09 12:11:22 INFO mapreduce.Job: map 100% reduce 100%  
19/09/09 12:11:24 INFO mapreduce.Job: Job job_1564775008355_5017 completed succe  
ssfully  
19/09/09 12:11:24 INFO mapreduce.Job: Counters: 49  
File System Counters  
FILE: Number of bytes read=69  
FILE: Number of bytes written=475039  
FILE: Number of read operations=0  
FILE: Number of large read operations=0  
FILE: Number of write operations=0  
HDFS: Number of bytes read=630  
HDFS: Number of bytes written=17  
HDFS: Number of read operations=9  
HDFS: Number of large read operations=0  
HDFS: Number of write operations=2  
Job Counters  
Launched map tasks=2  
Launched reduce tasks=1  
Data-local map tasks=2
```

```
rs6785@login-2-1:~  
Launched reduce tasks=1  
Data-local map tasks=2  
Total time spent by all maps in occupied slots (ms)=18036  
Total time spent by all reduces in occupied slots (ms)=13662  
Total time spent by all map tasks (ms)=4509  
Total time spent by all reduce tasks (ms)=2277  
Total vcore-milliseconds taken by all map tasks=4509  
Total vcore-milliseconds taken by all reduce tasks=2277  
Total megabyte-milliseconds taken by all map tasks=18468864  
Total megabyte-milliseconds taken by all reduce tasks=13989888  
Map-Reduce Framework  
Map input records=5  
Map output records=5  
Map output bytes=55  
Map output materialized bytes=92  
Input split bytes=210  
Combine input records=0  
Combine output records=0  
Reduce input groups=2  
Reduce shuffle bytes=92  
Reduce input records=5  
Reduce output records=2  
Spilled Records=10  
Shuffled Maps =2  
Failed Shuffles=0  
Merged Map outputs=2  
GC time elapsed (ms)=379  
CPU time spent (ms)=7870  
Physical memory (bytes) snapshot=2029522944  
Virtual memory (bytes) snapshot=11212648448  
Total committed heap usage (bytes)=3985113088  
Shuffle Errors  
BAD_ID=0  
CONNECTION=0  
IO_ERROR=0  
WRONG_LENGTH=0  
WRONG_MAP=0  
WRONG_REDUCE=0  
File Input Format Counters  
Bytes Read=420  
File Output Format Counters  
Bytes Written=17  
19/09/09 12:11:24 INFO streaming.StreamJob: Output directory: /user/rs6785/class  
1/output
```

## Screenshot that the program ran successfully

```
rs6785@login-2-1:~  
Total vooere-milliseconds taken by all reduce tasks=2277  
Total megabyte-milliseconds taken by all map tasks=18468864  
Total megabyte-milliseconds taken by all reduce tasks=13989888  
Map-Reduce Framework  
  Map input records=5  
  Map output records=5  
  Map output bytes=55  
  Map output materialized bytes=92  
  Input split bytes=210  
  Combine input records=0  
  Combine output records=0  
  Reduce input groups=2  
  Reduce shuffle bytes=92  
  Reduce input records=5  
  Reduce output records=2  
  Spilled Records=10  
  Shuffled Maps =2  
  Failed Shuffles=0  
  Merged Map outputs=2  
  GC time elapsed (ms)=379  
  CPU time spent (ms)=7870  
  Physical memory (bytes) snapshot=2029522944  
  Virtual memory (bytes) snapshot=11212648448  
  Total committed heap usage (bytes)=3985113088  
Shuffle Errors  
  BAD_ID=0  
  CONNECTION=0  
  IO_ERROR=0  
  WRONG_LENGTH=0  
  WRONG_MAP=0  
  WRONG_REDUCE=0  
File Input Format Counters  
  Bytes Read=420  
File Output Format Counters  
  Bytes Written=17  
19/09/09 12:11:24 INFO streaming.StreamJob: Output directory: /user/rs6785/class  
1/output  
[rs6785@login-2-1 ~]$ hdfs dfs -ls /user/rs6785/class1/output  
Found 2 items  
-rw-r--r--+ 3 rs6785 users      0 2019-09-09 12:11 /user/rs6785/class1/outp  
ut/_SUCCESS  
-rw-r--r--+ 3 rs6785 users     17 2019-09-09 12:11 /user/rs6785/class1/outp  
ut/part-00000  
[rs6785@login-2-1 ~]$
```

## Output

```
rs6785@login-2-1:~  
Map-Reduce Framework  
  Map input records=5  
  Map output records=5  
  Map output bytes=55  
  Map output materialized bytes=92  
  Input split bytes=210  
  Combine input records=0  
  Combine output records=0  
  Reduce input groups=2  
  Reduce shuffle bytes=92  
  Reduce input records=5  
  Reduce output records=2  
  Spilled Records=10  
  Shuffled Maps =2  
  Failed Shuffles=0  
  Merged Map outputs=2  
  GC time elapsed (ms)=379  
  CPU time spent (ms)=7870  
  Physical memory (bytes) snapshot=2029522944  
  Virtual memory (bytes) snapshot=11212648448  
  Total committed heap usage (bytes)=3985113088  
Shuffle Errors  
  BAD_ID=0  
  CONNECTION=0  
  IO_ERROR=0  
  WRONG_LENGTH=0  
  WRONG_MAP=0  
  WRONG_REDUCE=0  
File Input Format Counters  
  Bytes Read=420  
File Output Format Counters  
  Bytes Written=17  
19/09/09 12:11:24 INFO streaming.StreamJob: Output directory: /user/rs6785/class  
1/output  
[rs6785@login-2-1 ~]$ hdfs dfs -ls /user/rs6785/class1/output  
Found 2 items  
-rw-r--r--+ 3 rs6785 users      0 2019-09-09 12:11 /user/rs6785/class1/outp  
ut/_SUCCESS  
-rw-r--r--+ 3 rs6785 users     17 2019-09-09 12:11 /user/rs6785/class1/outp  
ut/part-00000  
[rs6785@login-2-1 ~]$ hdfs dfs -cat /user/rs6785/class1/output/part-00000  
1949    111  
1950    22  
[rs6785@login-2-1 ~]$
```