

Job Output

```
rs6785@login-1-1:~  
[rs6785@login-1-1 ~]$ hadoop jar /opt/cloudera/parcels/CDH/lib/hadoop-mapreduce/  
hadoop-streaming.jar -D mapreduce.job.reduces=1 -files hdfs://dumbo/user/rs6785/  
class2/python_code/StringCountMapper.py,hdfs://dumbo/user/rs6785/class2/python_c  
ode/StringCountReducer.py -mapper "StringCountMapper.py" -reducer "StringCountRe  
ducer.py" -input /user/rs6785/class2/SampleInput.txt -output /user/rs6785/class2  
/output  
packageJobJar: [] [/opt/cloudera/parcels/CDH-5.15.2-1.cdh5.15.2.p0.3/jars/hadoop  
-streaming-2.6.0-cdh5.15.2.jar] /tmp/streamjob6834463529048293190.jar tmpDir=nul  
l  
19/09/13 19:53:52 INFO mapreduce.JobSubmissionFiles: Permissions on staging dire  
ctory /user/rs6785/.staging are incorrect: rwxrwx---. Fixing permissions to corr  
ect value rwxrwxr-x  
19/09/13 19:53:52 INFO mapred.FileInputFormat: Total input paths to process : 1  
19/09/13 19:53:53 INFO mapreduce.JobSubmitter: number of splits:2  
19/09/13 19:53:53 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_15  
64775008355_5502  
19/09/13 19:53:53 INFO impl.YarnClientImpl: Submitted application application_15  
64775008355_5502  
19/09/13 19:53:53 INFO mapreduce.Job: The url to track the job: http://babar.es.  
its.nyu.edu:8088/proxy/application_1564775008355_5502/  
19/09/13 19:53:53 INFO mapreduce.Job: Running job: job_1564775008355_5502  
19/09/13 19:53:58 INFO mapreduce.Job: Job job_1564775008355_5502 running in uber  
mode : false  
19/09/13 19:53:58 INFO mapreduce.Job: map 0% reduce 0%  
19/09/13 19:54:02 INFO mapreduce.Job: map 50% reduce 0%  
19/09/13 19:54:03 INFO mapreduce.Job: map 100% reduce 0%  
19/09/13 19:54:08 INFO mapreduce.Job: map 100% reduce 100%  
19/09/13 19:54:09 INFO mapreduce.Job: Job job_1564775008355_5502 completed succe  
ssfully  
19/09/13 19:54:09 INFO mapreduce.Job: Counters: 49  
File System Counters  
FILE: Number of bytes read=89  
FILE: Number of bytes written=474952  
FILE: Number of read operations=0  
FILE: Number of large read operations=0  
FILE: Number of write operations=0  
HDFS: Number of bytes read=632  
HDFS: Number of bytes written=35  
HDFS: Number of read operations=9  
HDFS: Number of large read operations=0  
HDFS: Number of write operations=2  
Job Counters  
Launched map tasks=2
```

```
rs6785@login-1-1:~  
Job Counters  
Launched map tasks=2  
Launched reduce tasks=1  
Data-local map tasks=2  
Total time spent by all maps in occupied slots (ms)=18208  
Total time spent by all reduces in occupied slots (ms)=14568  
Total time spent by all map tasks (ms)=4552  
Total time spent by all reduce tasks (ms)=2428  
Total vcore-milliseconds taken by all map tasks=4552  
Total vcore-milliseconds taken by all reduce tasks=2428  
Total megabyte-milliseconds taken by all map tasks=18644992  
Total megabyte-milliseconds taken by all reduce tasks=14917632  
Map-Reduce Framework  
Map input records=3  
Map output records=12  
Map output bytes=105  
Map output materialized bytes=132  
Input split bytes=198  
Combine input records=0  
Combine output records=0  
Reduce input groups=4  
Reduce shuffle bytes=132  
Reduce input records=12  
Reduce output records=4  
Spilled Records=24  
Shuffled Maps =2  
Failed Shuffles=0  
Merged Map outputs=2  
GC time elapsed (ms)=349  
CPU time spent (ms)=6650  
Physical memory (bytes) snapshot=2014629888  
Virtual memory (bytes) snapshot=11202093056  
Total committed heap usage (bytes)=4000317440  
Shuffle Errors  
BAD_ID=0  
CONNECTION=0  
IO_ERROR=0  
WRONG_LENGTH=0  
WRONG_MAP=0  
WRONG_REDUCE=0  
File Input Format Counters  
Bytes Read=434  
File Output Format Counters  
Bytes Written=35
```

Screenshot that the program ran successfully and the output

```
rs6785@login-1-1:~  
    Input split bytes=198  
    Combine input records=0  
    Combine output records=0  
    Reduce input groups=4  
    Reduce shuffle bytes=132  
    Reduce input records=12  
    Reduce output records=4  
    Spilled Records=24  
    Shuffled Maps =2  
    Failed Shuffles=0  
    Merged Map outputs=2  
    GC time elapsed (ms)=349  
    CPU time spent (ms)=6650  
    Physical memory (bytes) snapshot=2014629888  
    Virtual memory (bytes) snapshot=11202093056  
    Total committed heap usage (bytes)=4000317440  
Shuffle Errors  
    BAD_ID=0  
    CONNECTION=0  
    IO_ERROR=0  
    WRONG_LENGTH=0  
    WRONG_MAP=0  
    WRONG_REDUCE=0  
File Input Format Counters  
    Bytes Read=434  
File Output Format Counters  
    Bytes Written=35  
19/09/13 19:54:09 INFO streaming.StreamJob: Output directory: /user/rs6785/class  
2/output  
[rs6785@login-1-1 ~]$ hdfs dfs -ls /user/rs6785/class2/output  
Found 2 items  
-rw-r--r--+ 3 rs6785 users      0 2019-09-13 19:54 /user/rs6785/class2/outp  
ut/_SUCCESS  
-rw-r--r--+ 3 rs6785 users      35 2019-09-13 19:54 /user/rs6785/class2/outp  
ut/part-00000  
[rs6785@login-1-1 ~]$ hdfs dfs -ls /user/rs6785/class2/output/part-00000  
-rw-r--r--+ 3 rs6785 users      35 2019-09-13 19:54 /user/rs6785/class2/outp  
ut/part-00000  
[rs6785@login-1-1 ~]$ hdfs dfs -cat /user/rs6785/class2/output/part-00000  
Chicago 1  
Dec      3  
Java     0  
hackathon      2  
[rs6785@login-1-1 ~]$
```